

Container Networking

Baohua Yang

2016-06-27

Outline

- Introduction
- Access Mechanisms
- Carrier Network Mechanisms
- Potential Solutions
- Conclusion

Introduction

- Two fundamental questions (since VMs)
 - How to access/bind to the under carrier networks
 - How to provide the carrier networks?
- More requirements with Container Cloud
 - Scalability
 - Isolation
 - Dynamics
 - Virtual address (IP, Port)
 - Service Discovery
 - QoS
 - Integration
 - And ...



Today, we try to answer the 2 basic questions.

1


How to access the network?

Access Mechanisms

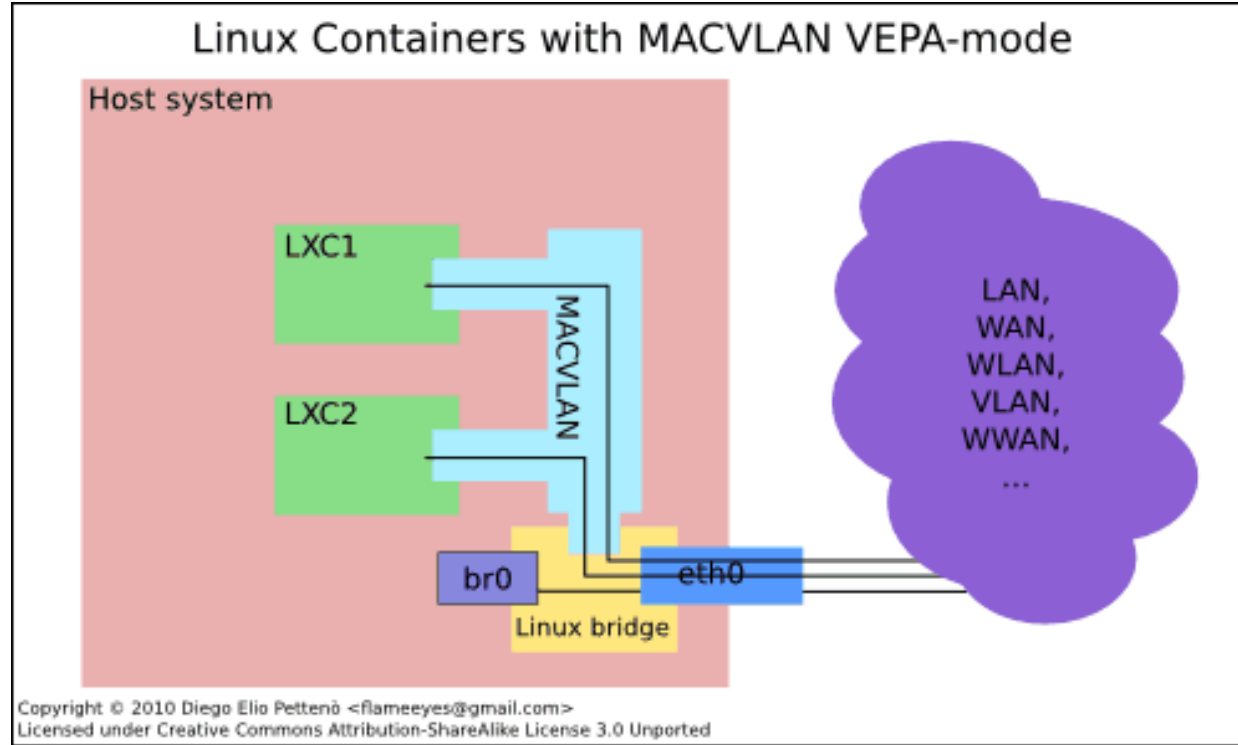
- Software Based
 - Virtual ports + Linux Bridges/ovs
- Hardware Based
 - Mapping virtual Nic
- Libnetwork
 - Supported since *1.7.0* (2015-06-16)
 - Docker official API to create network and bind container



Software Based

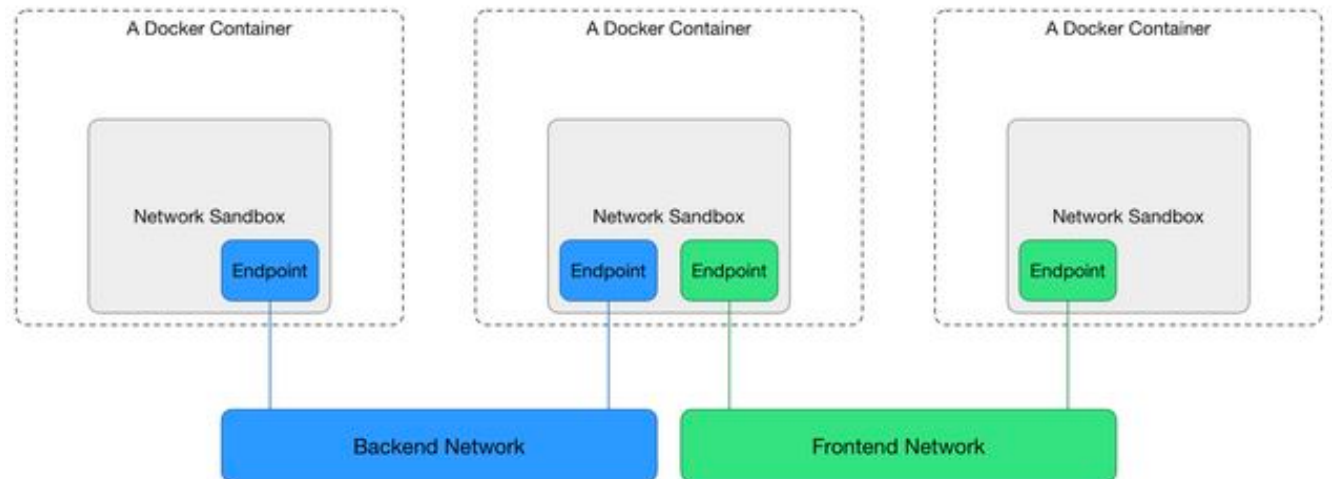


Hardware Based



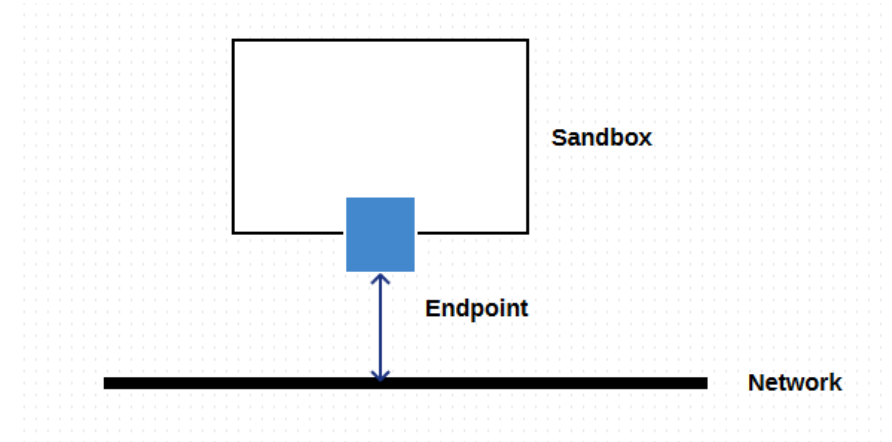
Libnetwork

- Experimentally support cross-host networking
- Based on socketplane
- Separated networking module from the core Docker engine
- New pluggable architecture
- Implements the Container Network Model (CNM)
- New network API/CLI



Libnetwork (cont.)

- Main Concepts:
 - Sandbox: Configuration of a container's network stack
 - Network: Endpoints that can communicate with Each other
 - Endpoint: Connects sandbox to networks
 - NetworkController: controller, exposes REST API to Docker Engine
 - Driver: handler real works (IPAM) for networkcontroller
- SDN style
 - Similar to Neutron Core concepts



Libnetwork (cont.)


- Supported drivers:
 - Built-in
 - Null: no network
 - Host: user host network stack
 - Bridge: traditional Docker networking (new implementation)
 - Overlay: multi-host networking
 - Remote: Connecting to Docker Network plugins
 - Uses JSON-RPC
 - Can utilize 3rd party networking solution

2


How to design the network?

Carrier Network Mechanisms

- NAT
- L2 switching
- L3 routing
- Overlay
 - VLAN
 - GRE
 - VXLAN
 - STT
 - NVGRE
 - GENEVE




NVGRE Encapsulation



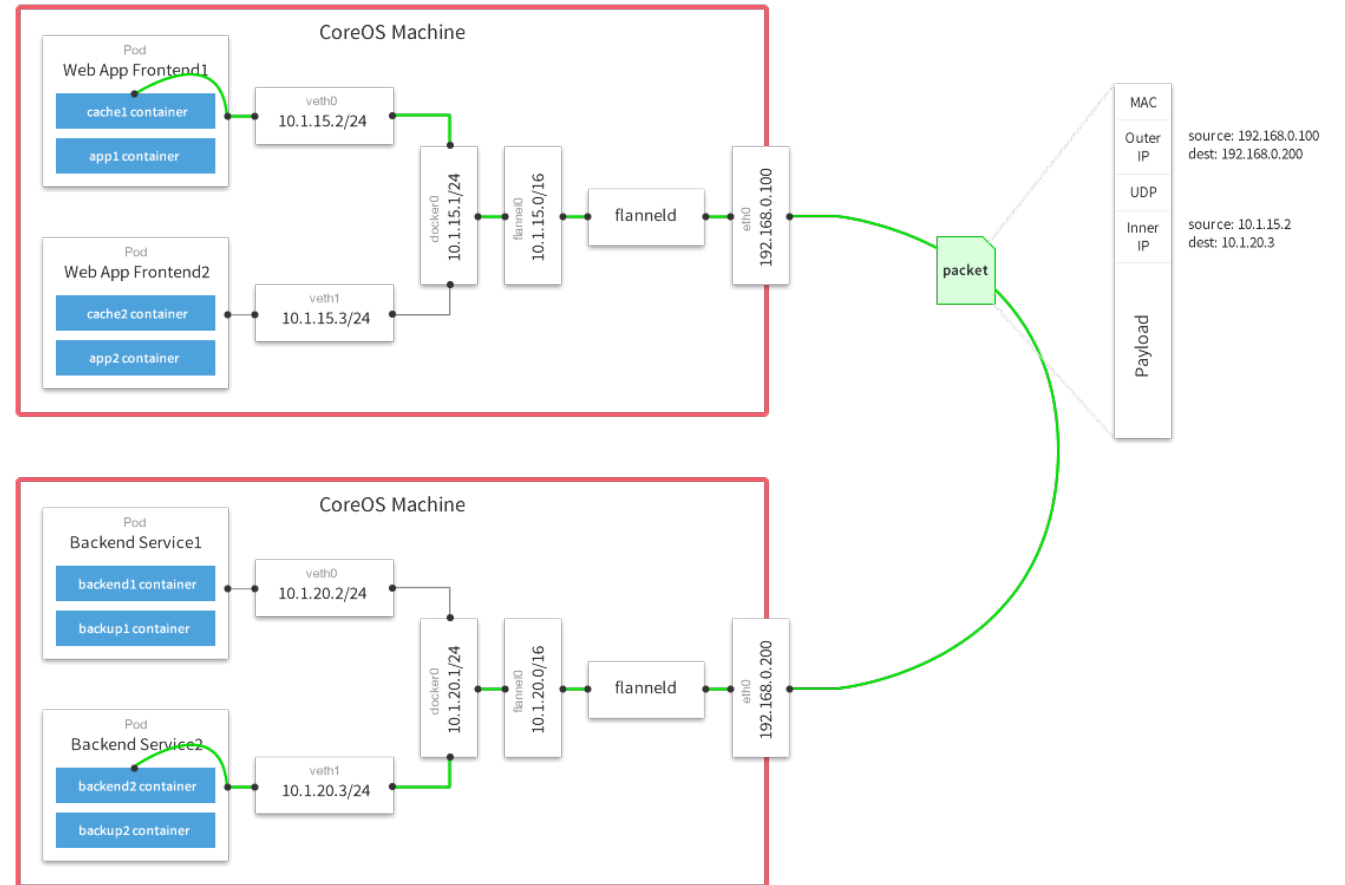
Potential Solutions

- NAT
- L2/L3 forwarding
- Overlay
- Flannel
- Weave
- Calico
- Neutron + Kuryr




Flannel

- CoreOS
- bridge + UDP Tunnel



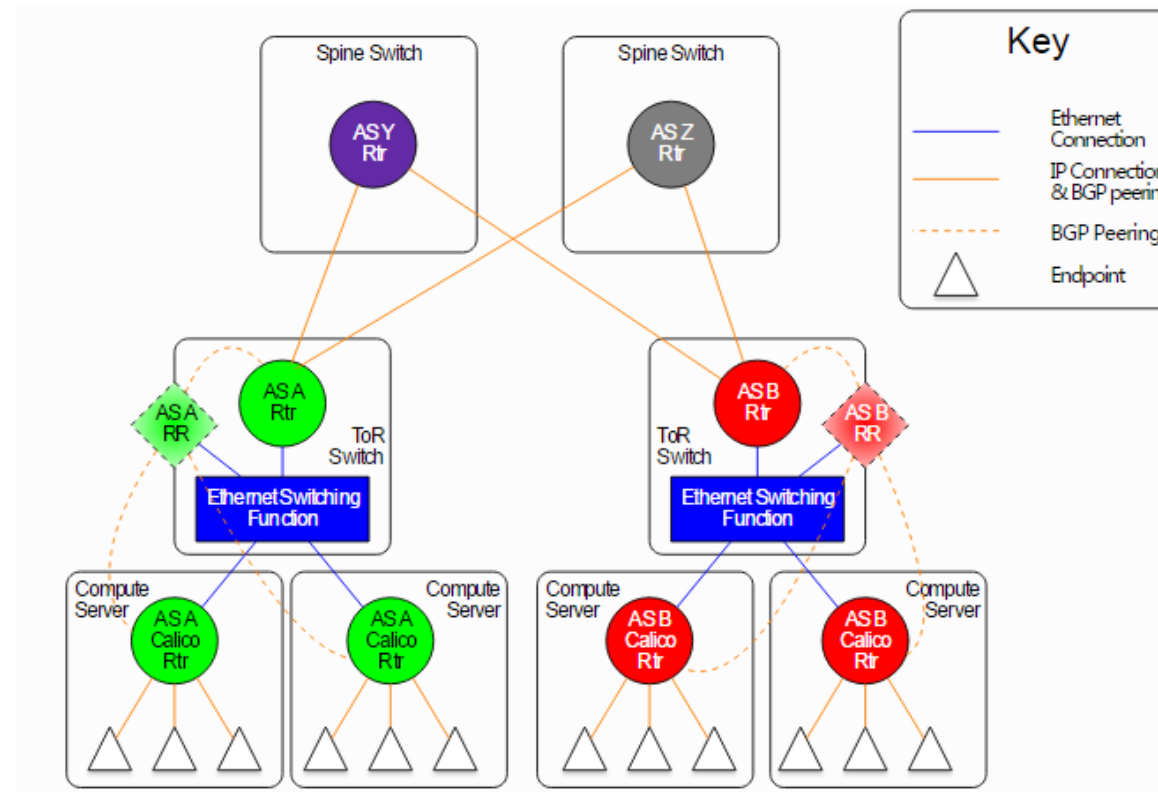
Weave

- Zett.io
- Pcap+UDP Tunnel




Calico

- Pure Layer 3 approach using vRouter + BGP
- Linux routing + iptables



Kuryr

- Bridge CNM with Openstack Neutron
- Kuryr is a Docker network plugin that uses Neutron
 - Provides networking services to Docker
 - Provides containerized images for the common Neutron plugins
 - Provides volume plugin for Docker (planned)



Kubernetes

- CNI
 - Rkt networking model
 - Container: a linux network namespace
 - Network: uniquely addressable entities that can communicate with each other
- Noop, bridge, local-host

Kuryr cont.

- Docker Networks: Neutron networks
- Docker Endpoints: Neutron ports
 - Neutron subnets gets created from a predefined Neutron subnetpool
 - Docker IPAM driver for Kuryr in the works
- Docker Join/Leave: plug/unplug

Kubernetes

- Abstractions

- pods:

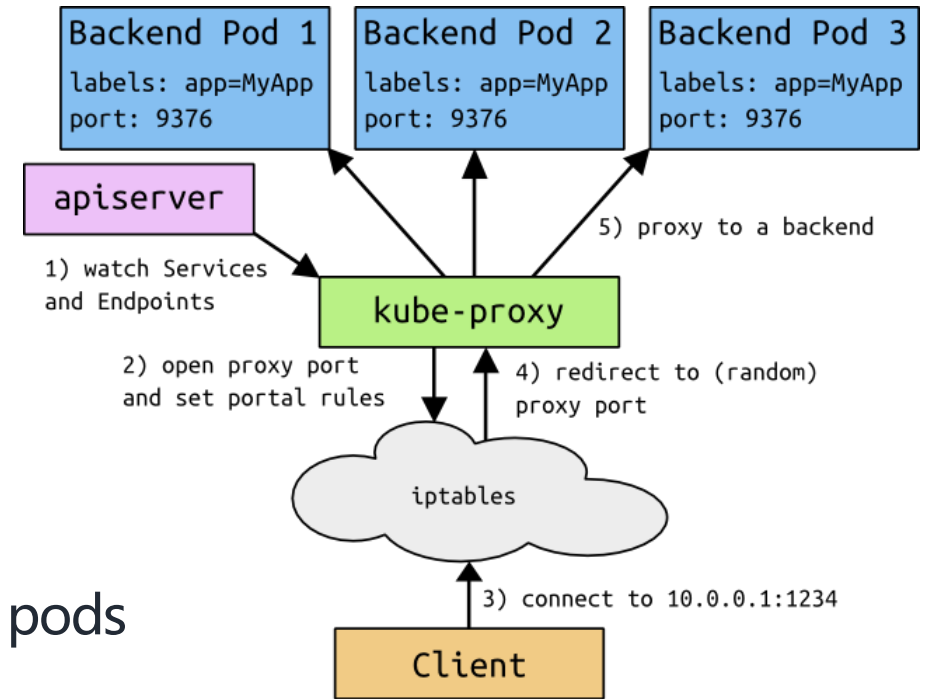
- group of containers on same host
 - IP per pod

- service:

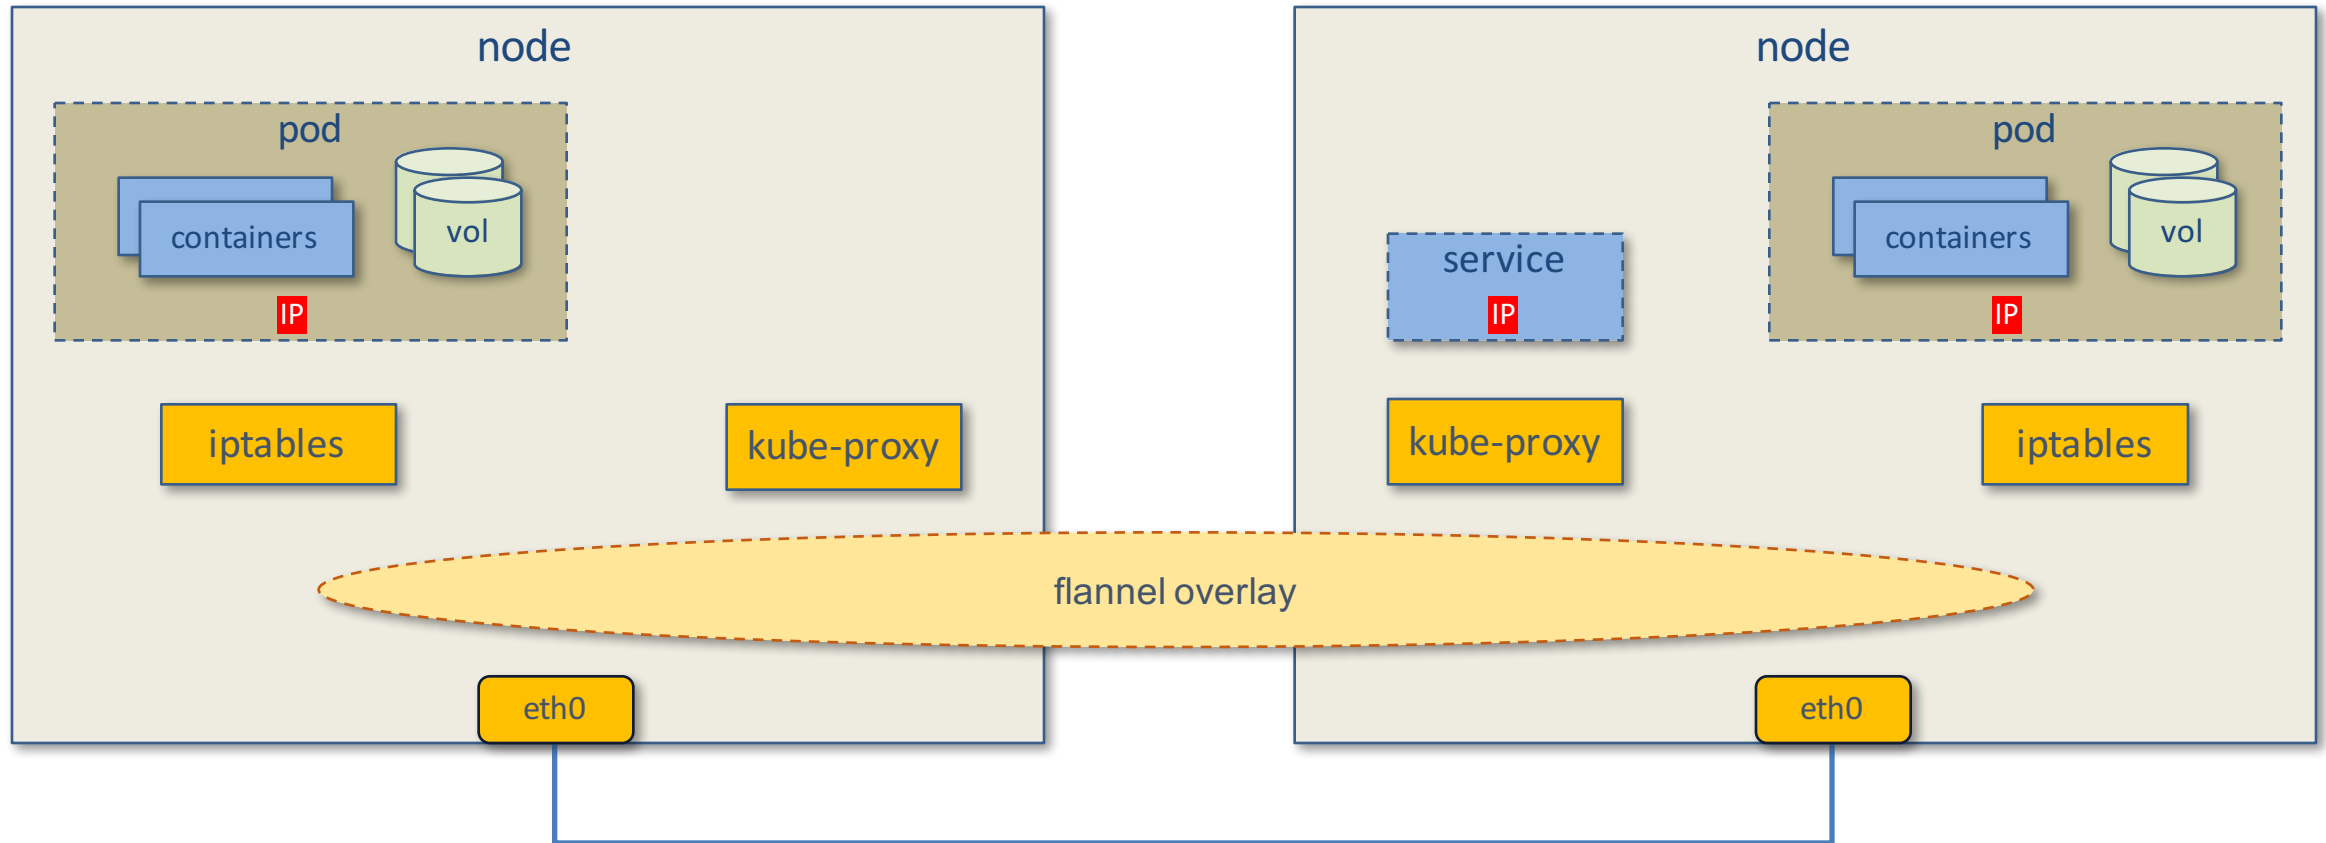
- proxy, load balancing
 - IP per service
 - replication controller: maintain exact number of pods

- Networking support

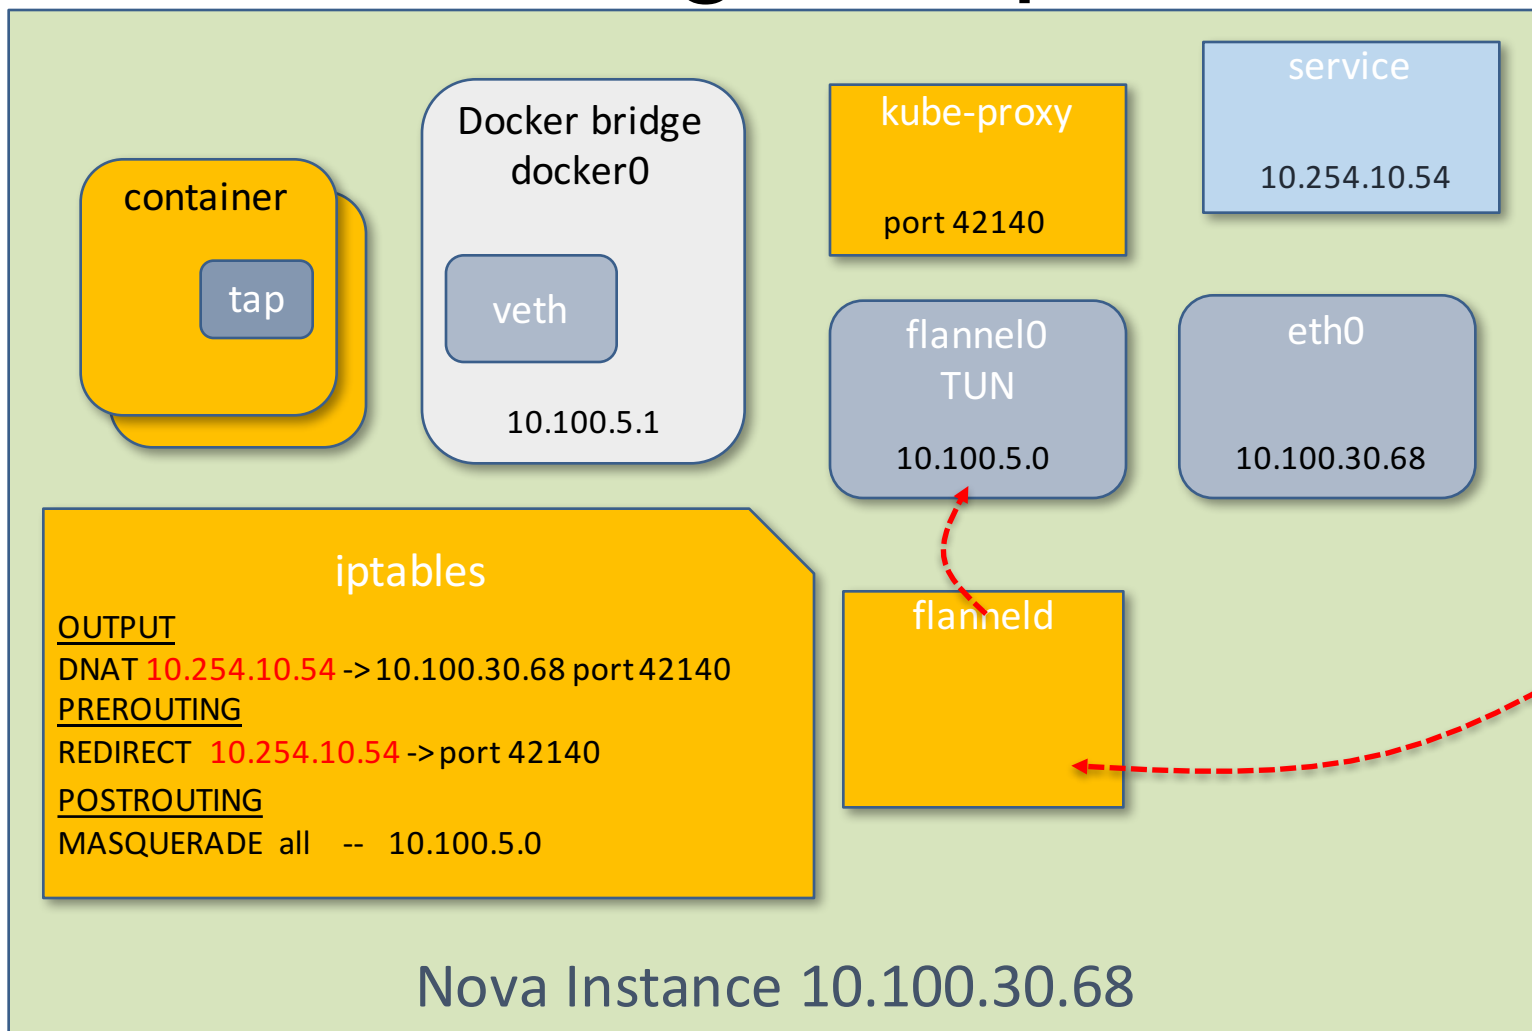
- kube-proxy: a Kubernetes component
 - flannel: an overlay network (other options available)
 - iptables rules: kernel support



Kubernetes Cluster in Operation



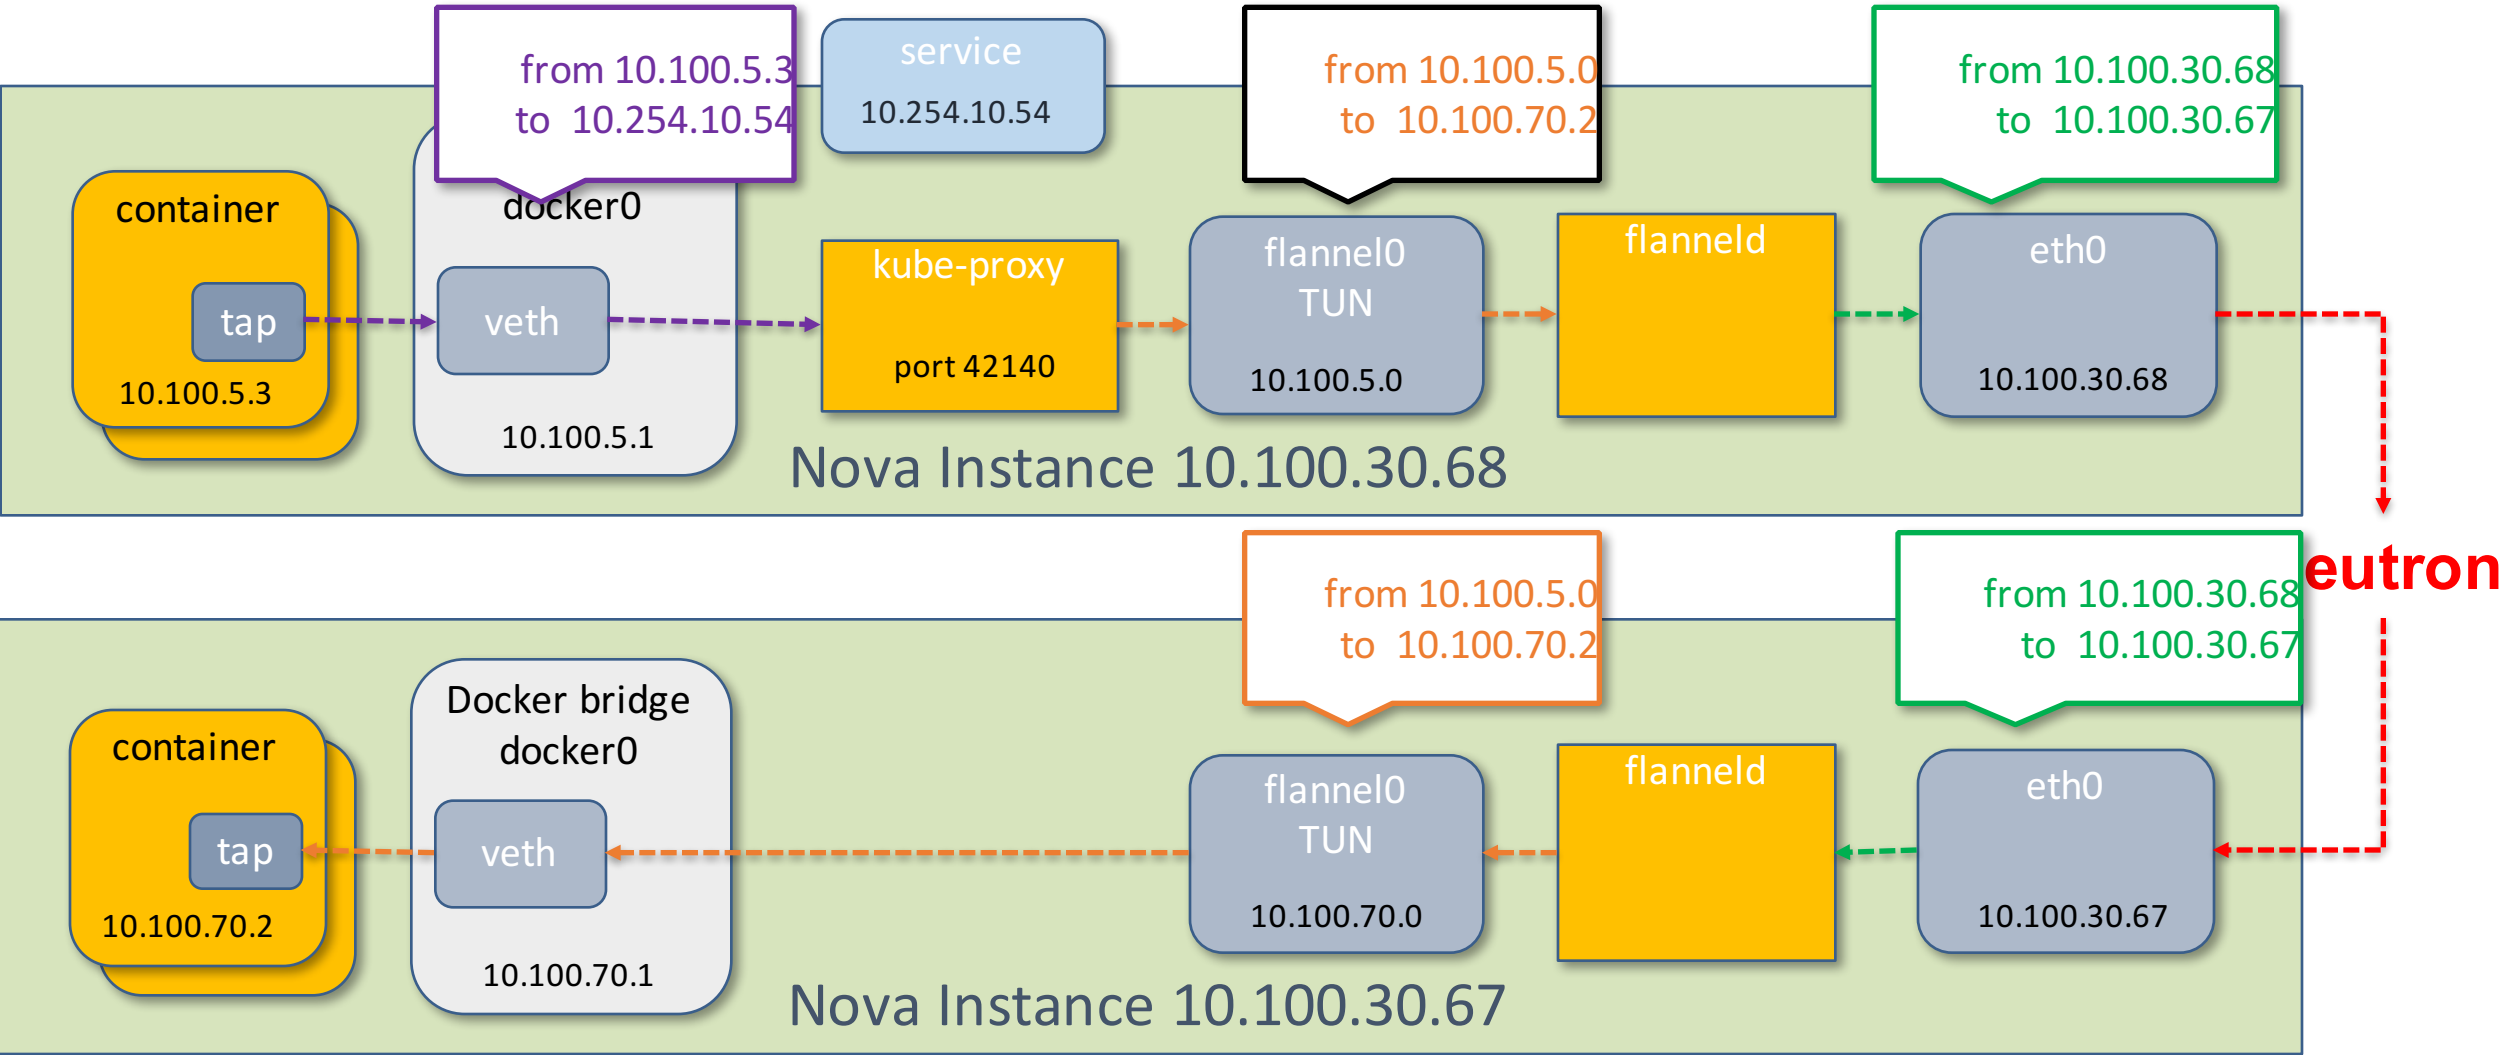
Networking Setup



```
JSON
├── action: "get"
├── node
│   ├── key: "/coreos.com/network/subnets"
│   ├── dir: true
│   └── nodes
│       ├── 0
│       │   ├── key: "/coreos.com/network/subnets/10.100.70.0-24"
│       │   ├── value: "{\"PublicIP\":\"10.100.30.67\"}"
│       │   ├── expiration: "2015-10-24T05:12:38.641417576Z"
│       │   ├── ttl: 50359
│       │   ├── modifiedIndex: 12
│       │   └── createdIndex: 12
│       └── 1
│           ├── key: "/coreos.com/network/subnets/10.100.5.0-24"
│           ├── value: "{\"PublicIP\":\"10.100.30.68\"}"
│           ├── expiration: "2015-10-24T05:12:44.454188039Z"
│           ├── ttl: 50364
│           ├── modifiedIndex: 13
│           └── createdIndex: 13
│   ├── modifiedIndex: 12
│   └── createdIndex: 12
```


etcd

Container-Container Communication



Neutron: VLAN-Aware VMs

- Provides an efficient way for interconnecting containers deployed within VMs
- Avoids using overlays on top of overlays
- Define new types of Neutron Ports
 - Trunk ports
 - Parent/Children relationship
- Initial patches under review
- Building momentum with increased interest



T: Trunk Port
S: Support

Optimistic Comparison

	Swam	Kubernetes	Mesos	OpenStack	CoreOS	CloudFoundry
NAT	Y	Y	Y	N	Y	Y
L2/L3 forwarding	Y	Y	Y	Y	Y	Y
Flannel	Y	Y	N	N	Y	N
Weave	Y	Y	Y	N	Y	N
Calico	Y	Y	Y	Y	Y	N
Neutron + Kuryr	Y	Y	Y	Y	N	N

* CF is discussing to support overlay.

Conclusion

- SDN is the de-facto standard in Cloud, everyone supports it!
- Libnetwork is key *in Docker ecosystem*.
- Host side networking is becoming more important.
- Many opportunities for new Cloud vendors.

Q&A

