

OpenStack Neutron

A system perspective

Baohua Yang
2014-04-11

Network as a Service



Neutron: Basic Concepts

- Minimal set of interfaces required for setting up networks for users

Network

isolated layer-2 broadcast domain;
private/shared

Subnet


CIDR IP address block associated with a `network`;
optionally associated gateway, DNS/DHCP servers

Port

virtual switch port on a `network`; has MAC and IP
address properties


A 3-tier App Example

- One possible implementation using a single router



Architecture overview

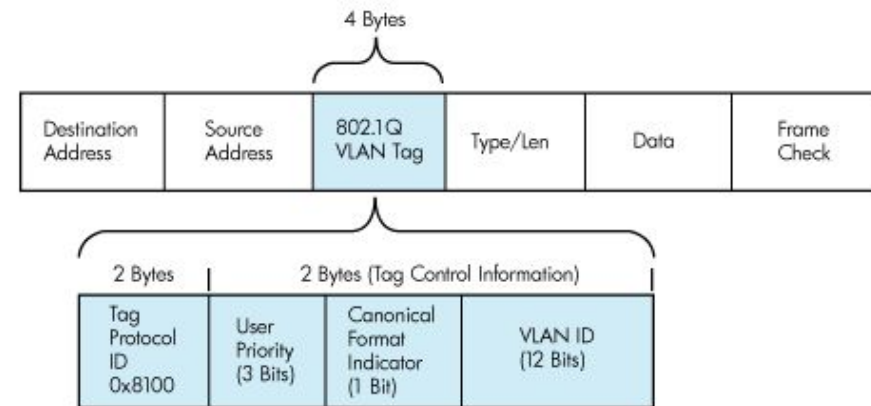
Big Picture



Prior Knowledge

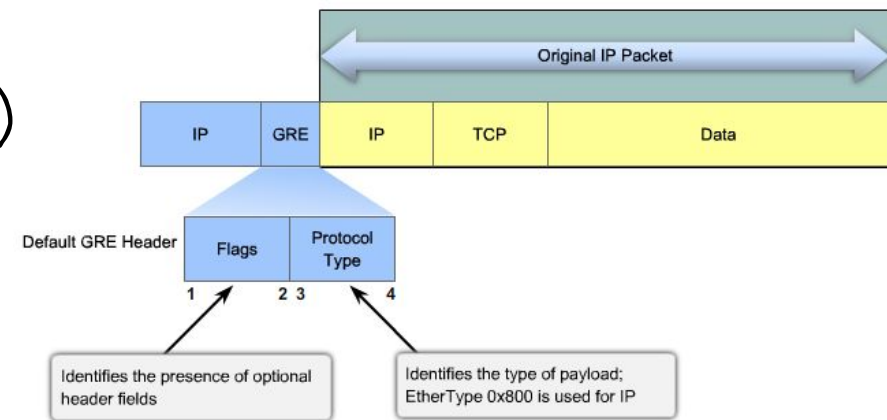
□ Vlan

- 802.1Q
- TPID : 16bit - 0x8100
- TCI : 16bit
 - PCP : 3bit
 - DEI : 1bit
 - VID : 12bit(0 ~ 4095)

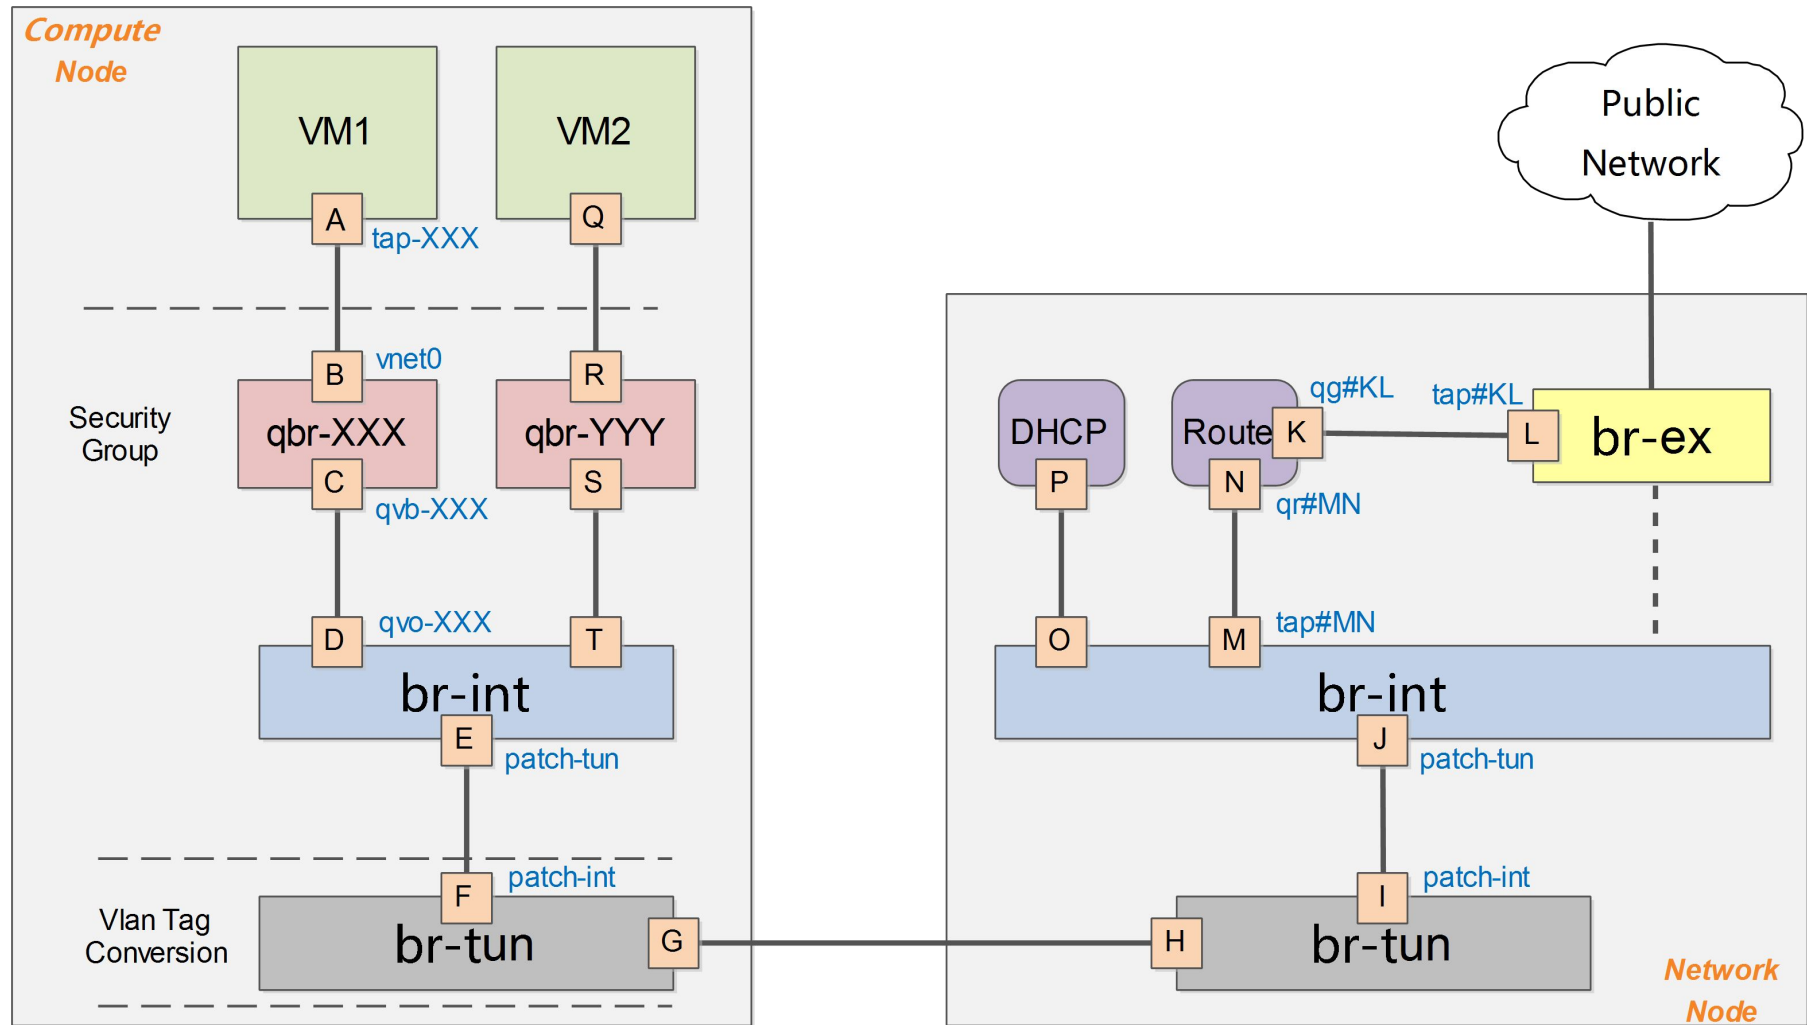


□ GRE

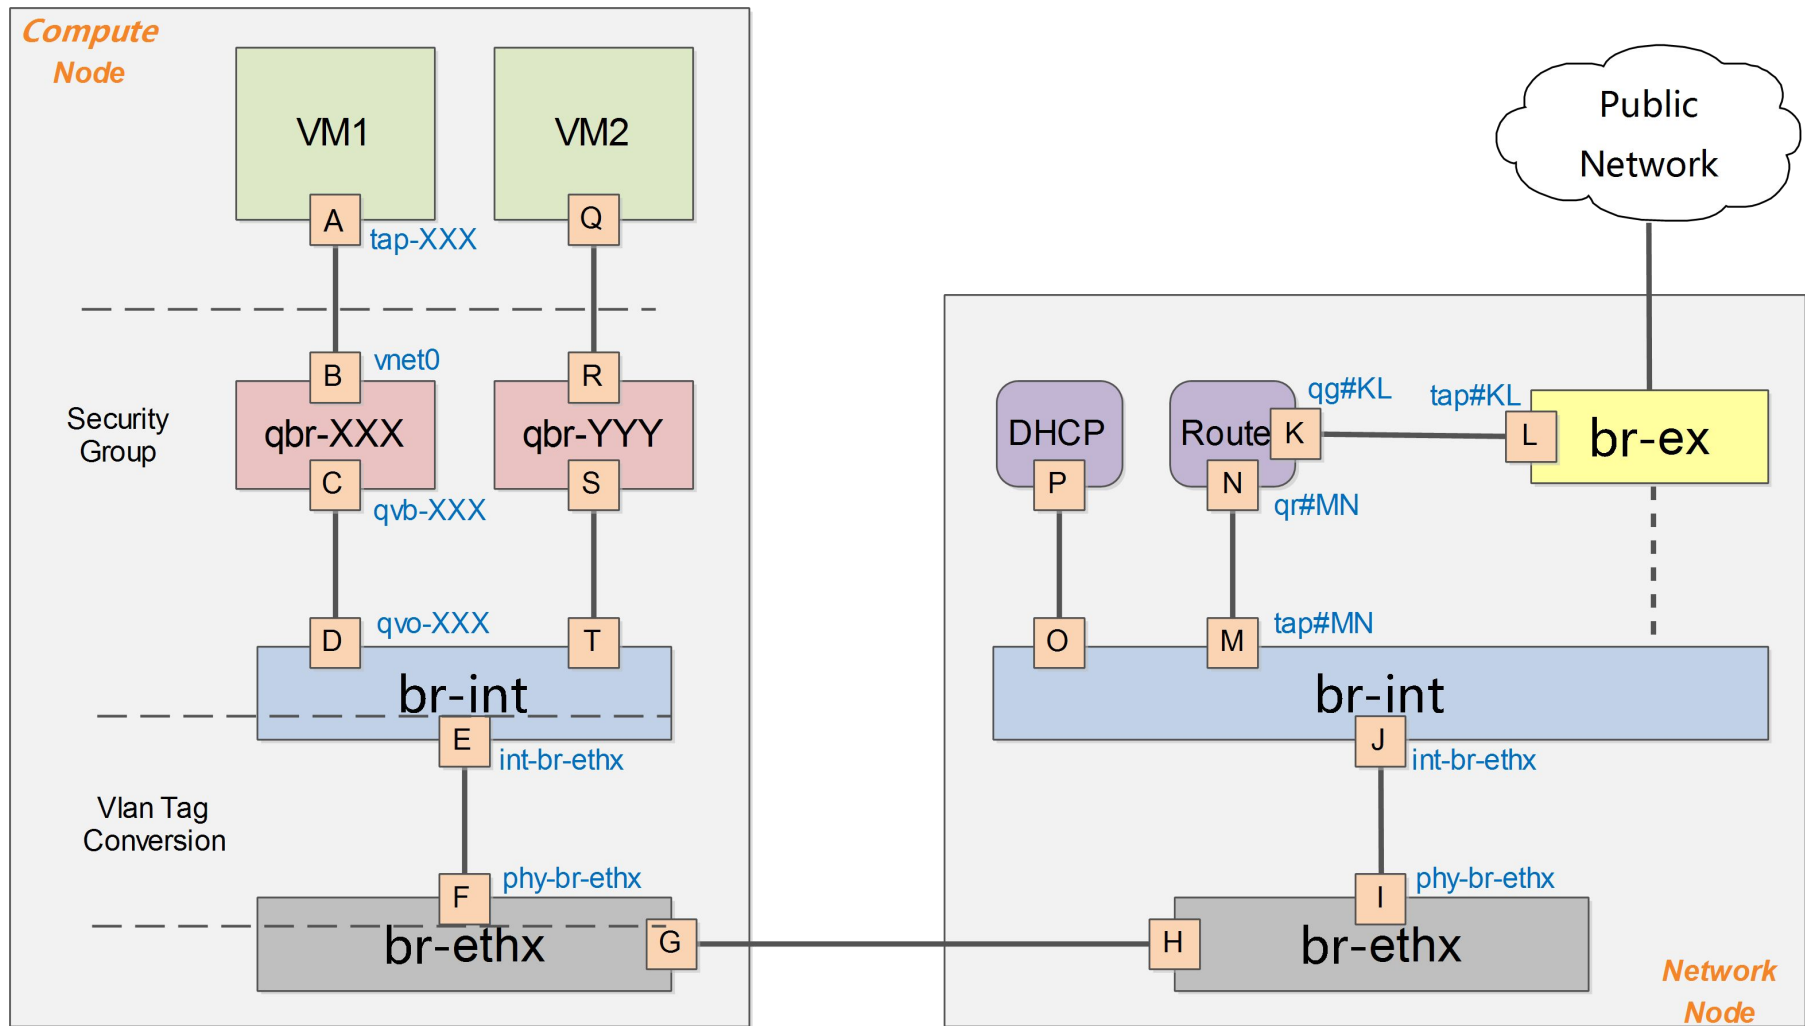
- 16 bytes header
- IP header



GRE Mode




Vlan Mode

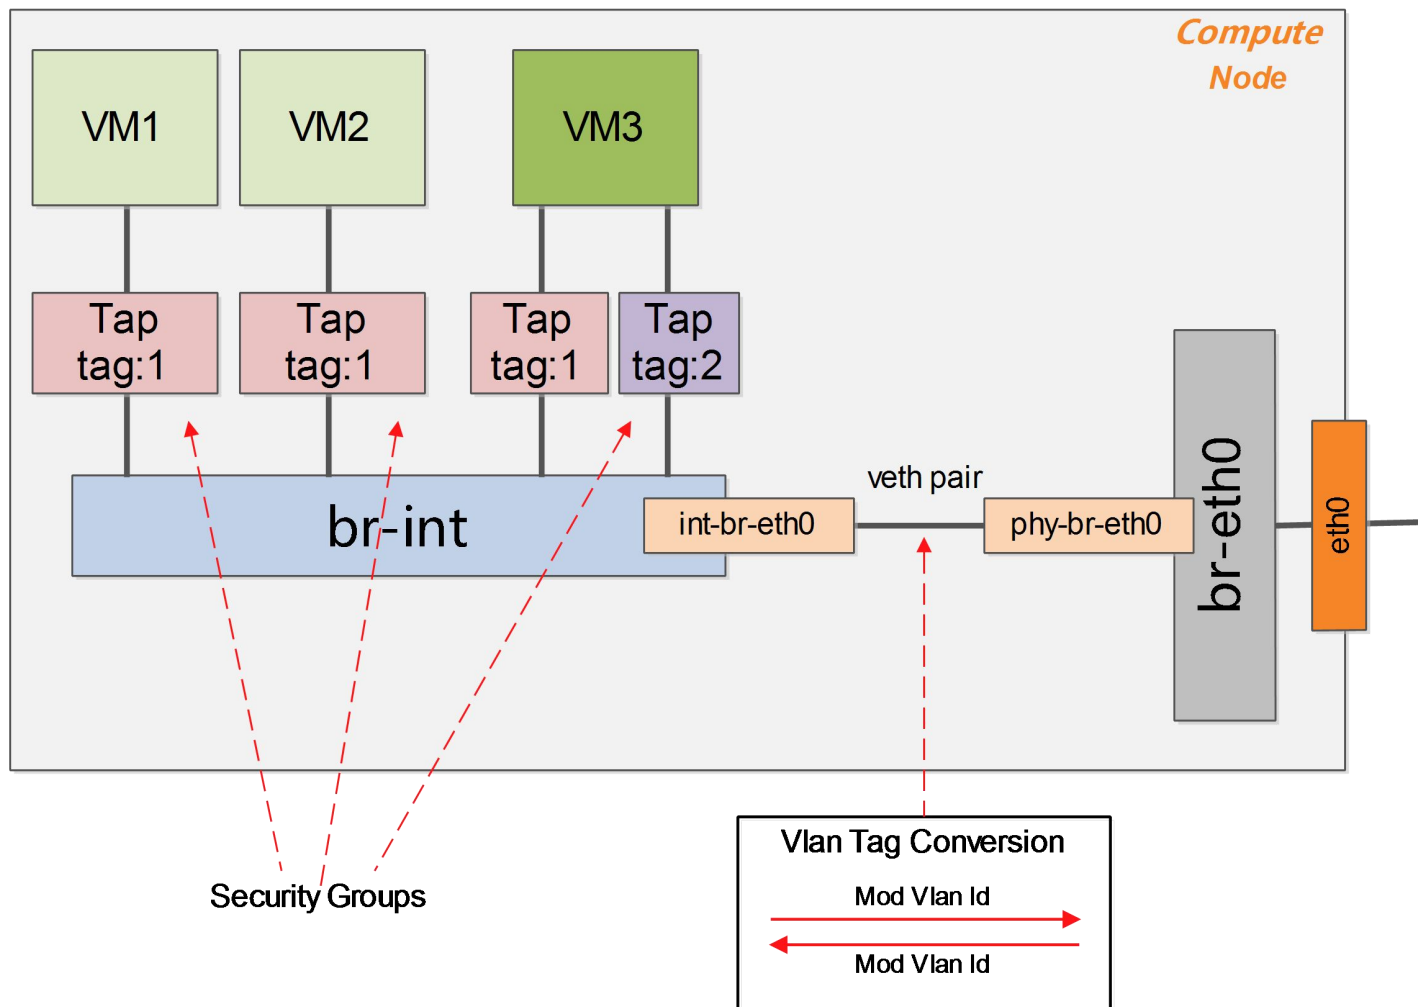


Walkthrough of Vlan Mode

Topology



Compute Node

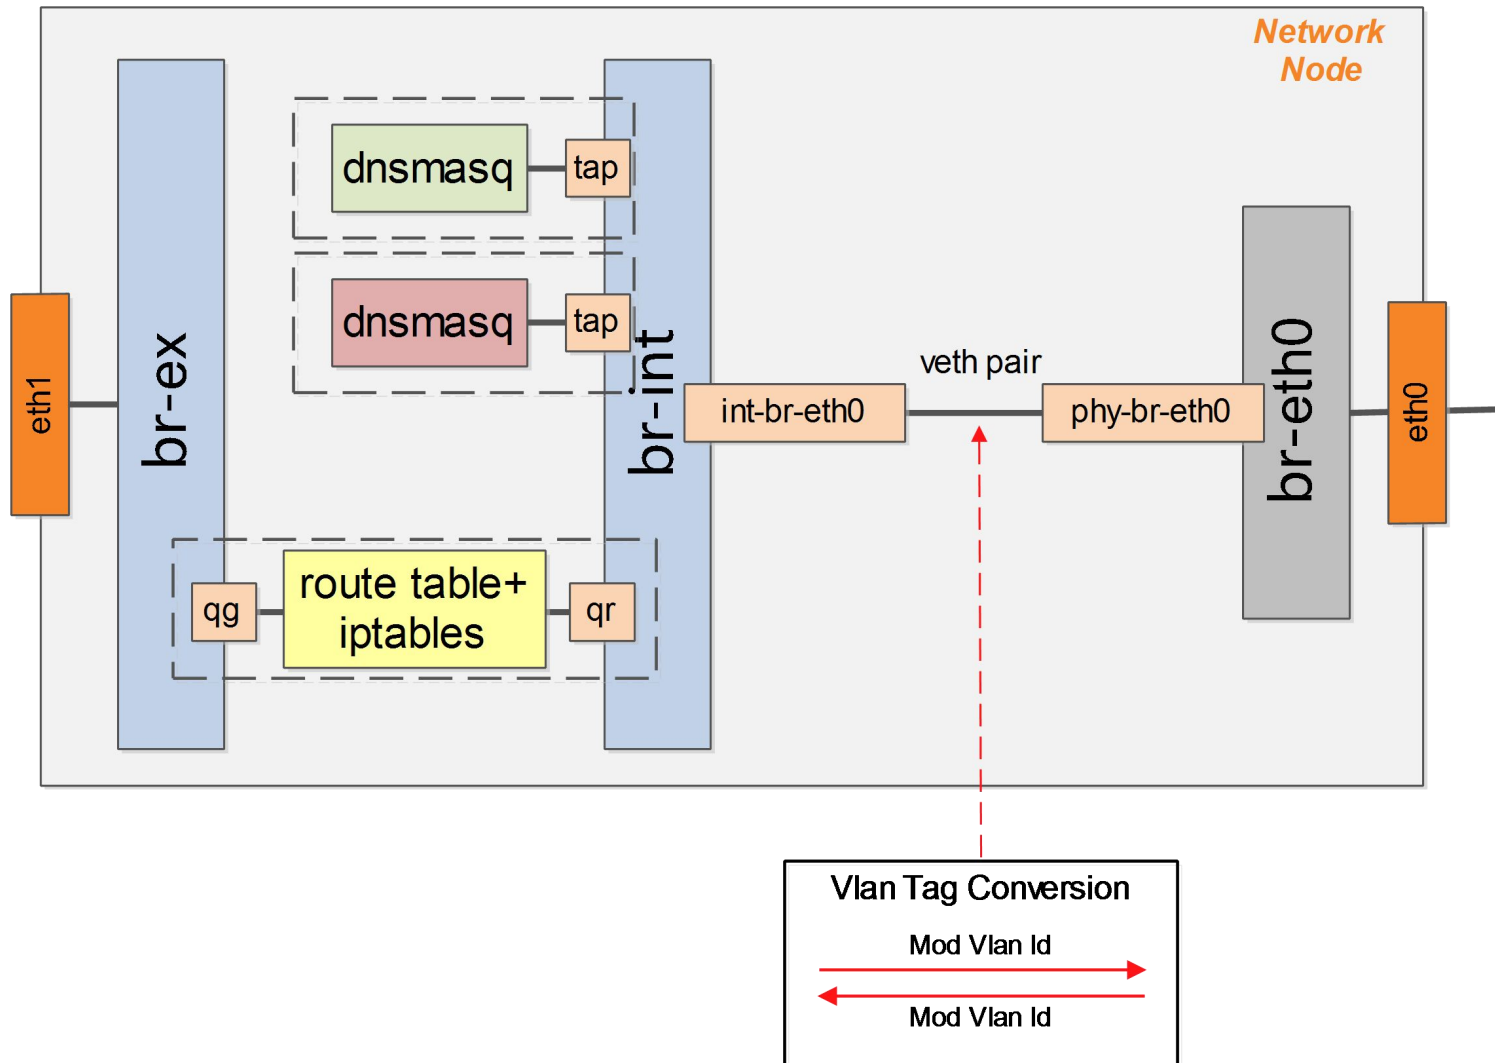


Compute Node

```
Easy0VS> list
br-eth0
Port:          br-eth0 phy-br-eth0 eth0
br-int
Port:          qvo260209fa-72 int-br-eth0 qvo8bf9cba2-3f qvod4de9fe0-6d br-int qvo583c7038-d3
Easy0VS> show br-int
Intf          Port      Vlan    Type      vmIP      vmMAC
int-br-eth0   20
qvo260209fa-72 11      1       192.168.0.4 fa:16:3e:0f:17:04
qvo583c7038-d3 2       1       192.168.0.2 fa:16:3e:9c:dc:3a
qvo8bf9cba2-3f 9       1       192.168.0.5 fa:16:3e:a2:2f:0e
qvod4de9fe0-6d 8       2       10.0.0.2     fa:16:3e:38:2b:2e
br-int        LOCAL      internal
Easy0VS> dump br-int
ID TAB PKT      PRI  MATCH                                ACT
0 0 6324      3    in=20,vlan=3                        mod_vlan_vid:2,NORMAL
1 0 17965     3    in=20,vlan=1                        mod_vlan_vid:1,NORMAL
2 0 6         2    in=20                                drop
3 0 34011     1    *                                    NORMAL
Easy0VS>

Easy0VS> show br-eth0
Intf          Port      Vlan    Type
eth0          1
phy-br-eth0   14
br-eth0        LOCAL      internal
Easy0VS> dump br-eth0
ID TAB PKT      PRI  MATCH                                ACT
0 0 28677     4    in=14,vlan=1                        mod_vlan_vid:1,NORMAL
1 0 6697      4    in=14,vlan=2                        mod_vlan_vid:3,NORMAL
2 0 9         2    in=14                                drop
3 0 25255     1    *                                    NORMAL
```

Network Node



Network Node

```
EasyOVS> list
br-eth0
  Port:          br-eth0 phy-br-eth0 eth0
br-ex
  Port:          br-ex eth1 qg-9b2db4ac-31
br-int
  Port:          int-br-eth0 br-int qr-2a169bb4-4d tapb66fe81c-de tapdb2f5a49-7c
EasyOVS> show br-int
Intf          Port      Vlan    Type      vmIP      vmMAC
int-br-eth0   8
br-int        LOCAL      internal
qr-2a169bb4-4d 2        2      internal  192.168.0.1  fa:16:3e:2f:e9:72
tapb66fe81c-de 4        1      internal  10.0.0.3    fa:16:3e:38:7d:3d
tapdb2f5a49-7c 3        2      internal  192.168.0.3  fa:16:3e:17:5c:36
EasyOVS> dump br-int
ID TAB PKT      PRI  MATCH                                     ACT
0  0  12          3    in=8,vlan=3                             mod_vlan_vid:1,NORMAL
1  0  41          3    in=8,vlan=1                             mod_vlan_vid:2,NORMAL
2  0  12          2    in=8                                     drop
3  0  44          1    *                                       NORMAL
EasyOVS> show br-eth0
Intf          Port      Vlan    Type
eth0          1
phy-br-eth0   6
br-eth0       LOCAL      internal
EasyOVS> dump br-eth0
ID TAB PKT      PRI  MATCH                                     ACT
0  0  18          4    in=6,vlan=1                             mod_vlan_vid:3,NORMAL
1  0  48          4    in=6,vlan=2                             mod_vlan_vid:1,NORMAL
2  0  6           2    in=6                                     drop
3  0  105         1    *                                       NORMAL
```

Advanced Topics

- Network Namespace
- Floating IP
- Security Group
- VXLAN
- ML2
- Multihost

Network Namespace

- Network namespace isolates the *network interface controllers* (physical or virtual), *iptables firewall rules*, *routing tables* etc.
- Network namespaces can be connected with each other using the "veth" virtual Ethernet device.

```
ip netns list  
ip netns add new_ns  
ip link add veth0 type veth peer name veth1  
ip link set veth1 netns new_ns  
ip netns exec new_ns <commands>
```

Network Namespace

```
root@Control:~#ip netns list
qdhcp-39edbf9b-a6da-4bab-8500-39ad91ed1984
qdhcp-035179eb-9022-4656-b88a-8bc841034eda
qrouter-03266ec4-a03b-41b2-897b-c18ae3279933
root@Control:~#ip netns exec qrouter-03266ec4-a03b-41b2-897b-c18ae3279933 ip addr
12: lo: <LOOPBACK,UP,LOWER_UP> mtu 16436 qdisc noqueue state UNKNOWN
    link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
    inet 127.0.0.1/8 scope host lo
    inet6 ::1/128 scope host
        valid_lft forever preferred_lft forever
15: qr-2a169bb4-4d: <BROADCAST,UP,LOWER_UP> mtu 1500 qdisc noqueue state UNKNOWN
    link/ether fa:16:3e:2f:e9:72 brd ff:ff:ff:ff:ff:ff
    inet 192.168.0.1/24 brd 192.168.0.255 scope global qr-2a169bb4-4d
    inet6 fe80::f816:3eff:fe2f:e972/64 scope link
        valid_lft forever preferred_lft forever
16: qg-9b2db4ac-31: <BROADCAST,UP,LOWER_UP> mtu 1500 qdisc noqueue state UNKNOWN
    link/ether fa:16:3e:4e:f1:b5 brd ff:ff:ff:ff:ff:ff
    inet 192.168.122.200/24 brd 192.168.122.255 scope global qg-9b2db4ac-31
    inet 192.168.122.201/32 brd 192.168.122.201 scope global qg-9b2db4ac-31
    inet 192.168.122.203/32 brd 192.168.122.203 scope global qg-9b2db4ac-31
    inet6 fe80::f816:3eff:fe4e:f1b5/64 scope link
        valid_lft forever preferred_lft forever
```

Floating IP

□ Route table + NAT

```
root@Control:~#ip netns exec qrouter-03266ec4-a03b-41b2-897b-c18ae3279933 ip route
192.168.0.0/24 dev qr-2a169bb4-4d proto kernel scope link src 192.168.0.1
192.168.122.0/24 dev qg-9b2db4ac-31 proto kernel scope link src 192.168.122.200
default via 192.168.122.1 dev qg-9b2db4ac-31
root@Control:~#ip netns exec qrouter-03266ec4-a03b-41b2-897b-c18ae3279933 iptables -t nat -S
-P PREROUTING ACCEPT
-P POSTROUTING ACCEPT
-P OUTPUT ACCEPT
-N neutron-l3-agent-OUTPUT
-N neutron-l3-agent-POSTROUTING
-N neutron-l3-agent-PREROUTING
-N neutron-l3-agent-float-snat
-N neutron-l3-agent-snat
-N neutron-postrouting-bottom
-A PREROUTING -j neutron-l3-agent-PREROUTING
-A POSTROUTING -j neutron-l3-agent-POSTROUTING
-A POSTROUTING -j neutron-postrouting-bottom
-A OUTPUT -j neutron-l3-agent-OUTPUT
-A neutron-l3-agent-OUTPUT -d 192.168.122.201/32 -j DNAT --to-destination 192.168.0.2
-A neutron-l3-agent-OUTPUT -d 192.168.122.203/32 -j DNAT --to-destination 192.168.0.5
-A neutron-l3-agent-POSTROUTING ! -i qg-9b2db4ac-31 ! -o qg-9b2db4ac-31 -m conntrack ! --ctstate DNAT -j ACCEPT
-A neutron-l3-agent-PREROUTING -d 169.254.169.254/32 -p tcp -m tcp --dport 80 -j REDIRECT --to-ports 9697
-A neutron-l3-agent-PREROUTING -d 192.168.122.201/32 -j DNAT --to-destination 192.168.0.2
-A neutron-l3-agent-PREROUTING -d 192.168.122.203/32 -j DNAT --to-destination 192.168.0.5
-A neutron-l3-agent-float-snat -s 192.168.0.2/32 -j SNAT --to-source 192.168.122.201
-A neutron-l3-agent-float-snat -s 192.168.0.5/32 -j SNAT --to-source 192.168.122.203
-A neutron-l3-agent-snat -j neutron-l3-agent-float-snat
-A neutron-l3-agent-snat -s 192.168.0.0/24 -j SNAT --to-source 192.168.122.200
-A neutron-postrouting-bottom -j neutron-l3-agent-snat
```


Security Group

```
root@Compute:~# easyovs -m "br-int show"
```

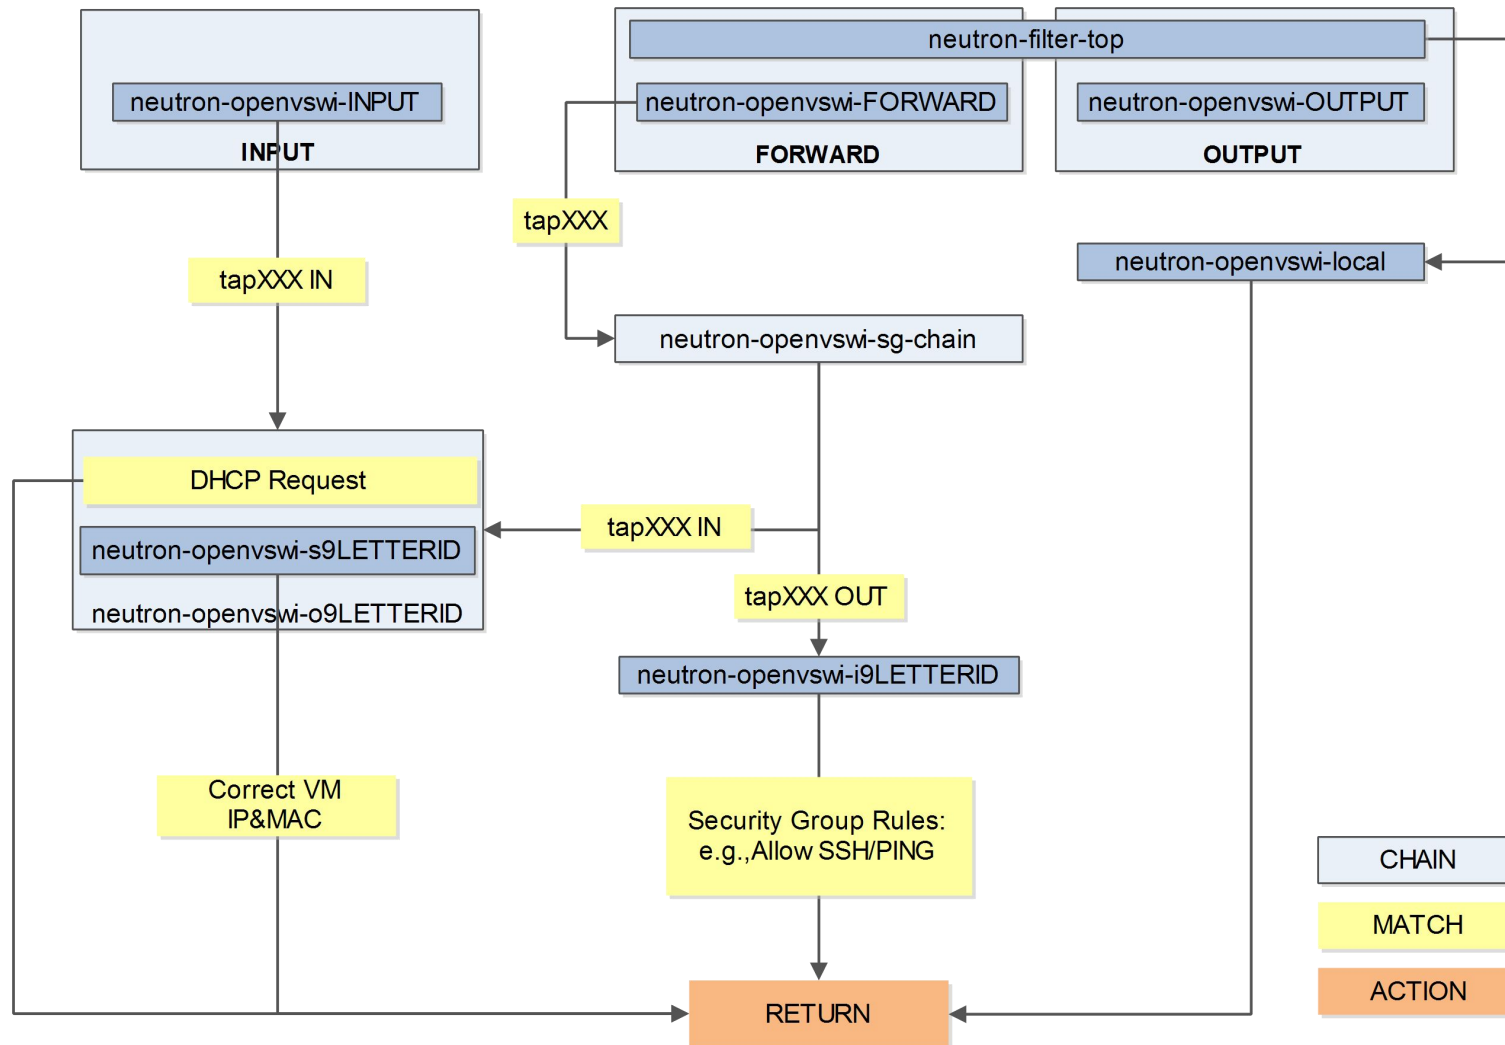
Intf	Port	Vlan	Type	vmIP	vmMAC
int-br-eth0	20				
qvo260209fa-72	11	1		192.168.0.4	fa:16:3e:0f:17:04
qvo583c7038-d3	2	1		192.168.0.2	fa:16:3e:9c:dc:3a
qvo8bf9cba2-3f	9	1		192.168.0.5	fa:16:3e:a2:2f:0e
qvod4de9fe0-6d	8	2		10.0.0.2	fa:16:3e:38:2b:2e
br-int	LOCAL		internal		

```
Easy0VS> ipt 192.168.0.2
```

```
## IP = 192.168.0.2, port = qvo583c7038-d ##
```


PKTS	SOURCE	DESTINATION	PROT	OTHER
#IN:				
672	all	all	all	state RELATED,ESTABLISHED
0	all	all	tcp	tcp dpt:22
0	all	all	icmp	
0	192.168.0.4	all	all	
3	192.168.0.5	all	all	
8	10.0.0.2	all	all	
85778	192.168.0.3	all	udp	udp spt:67 dpt:68
#OUT:				
196K	all	all	udp	udp spt:68 dpt:67
86149	all	all	all	state RELATED,ESTABLISHED
1241	all	all	all	
#SRC FILTER:				
59157	192.168.0.2	all	all	MAC FA:16:3E:9C:DC:3A

Security Group

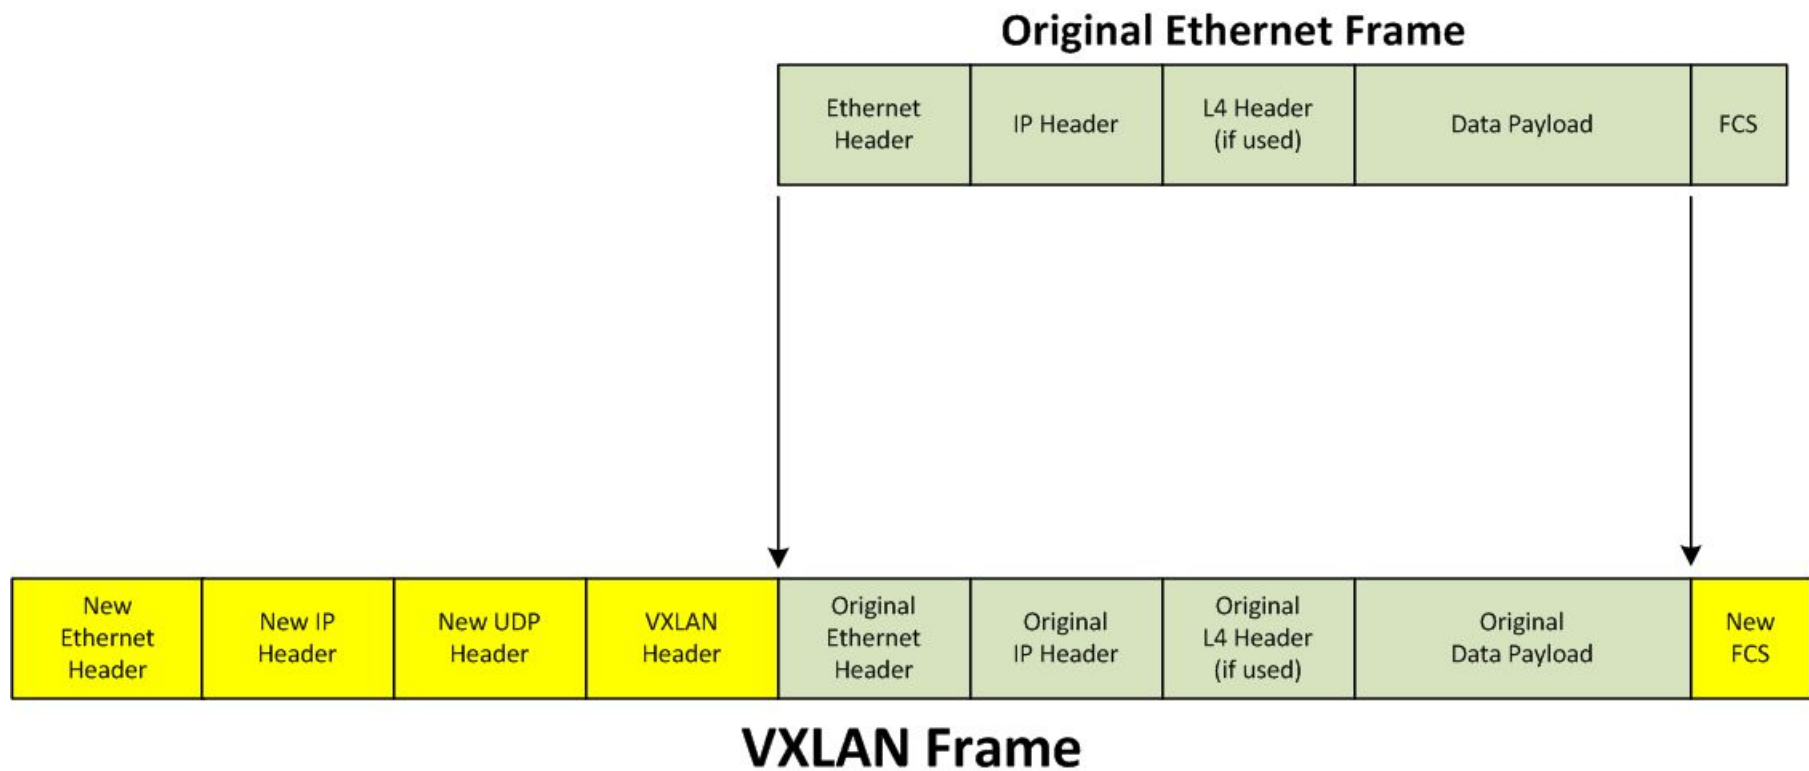


VXLAN

- Standardized overlay technology for encapsulating layer 2 traffic on top of an IP fabric




VXLAN



ML2

□ Before Modular Layer 2

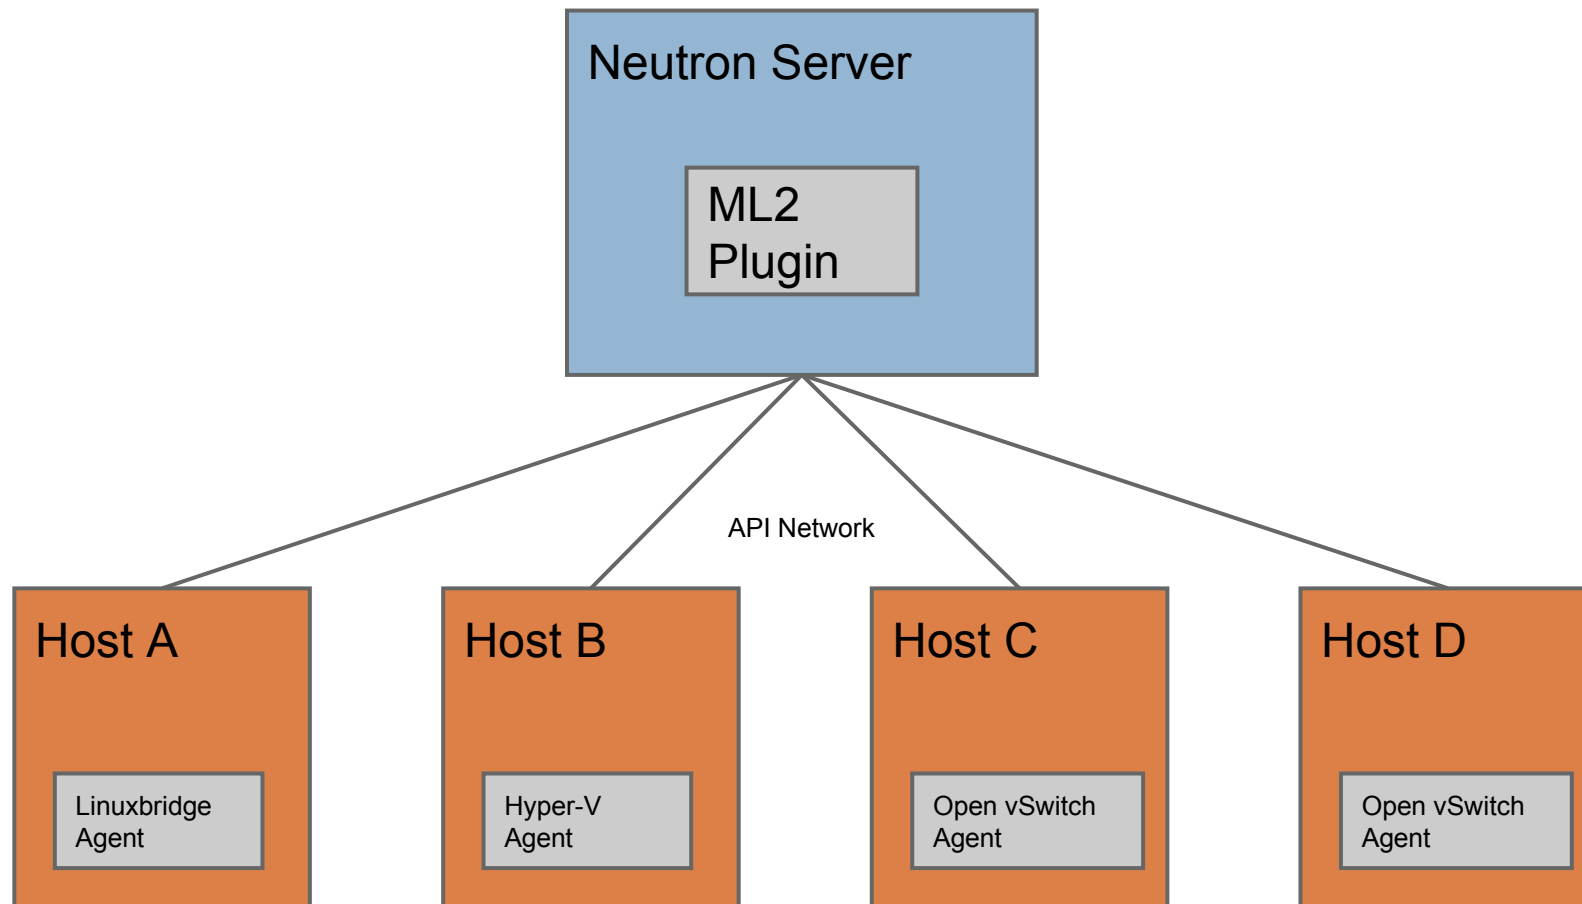


I want to
write a
Neutron
Plugin.

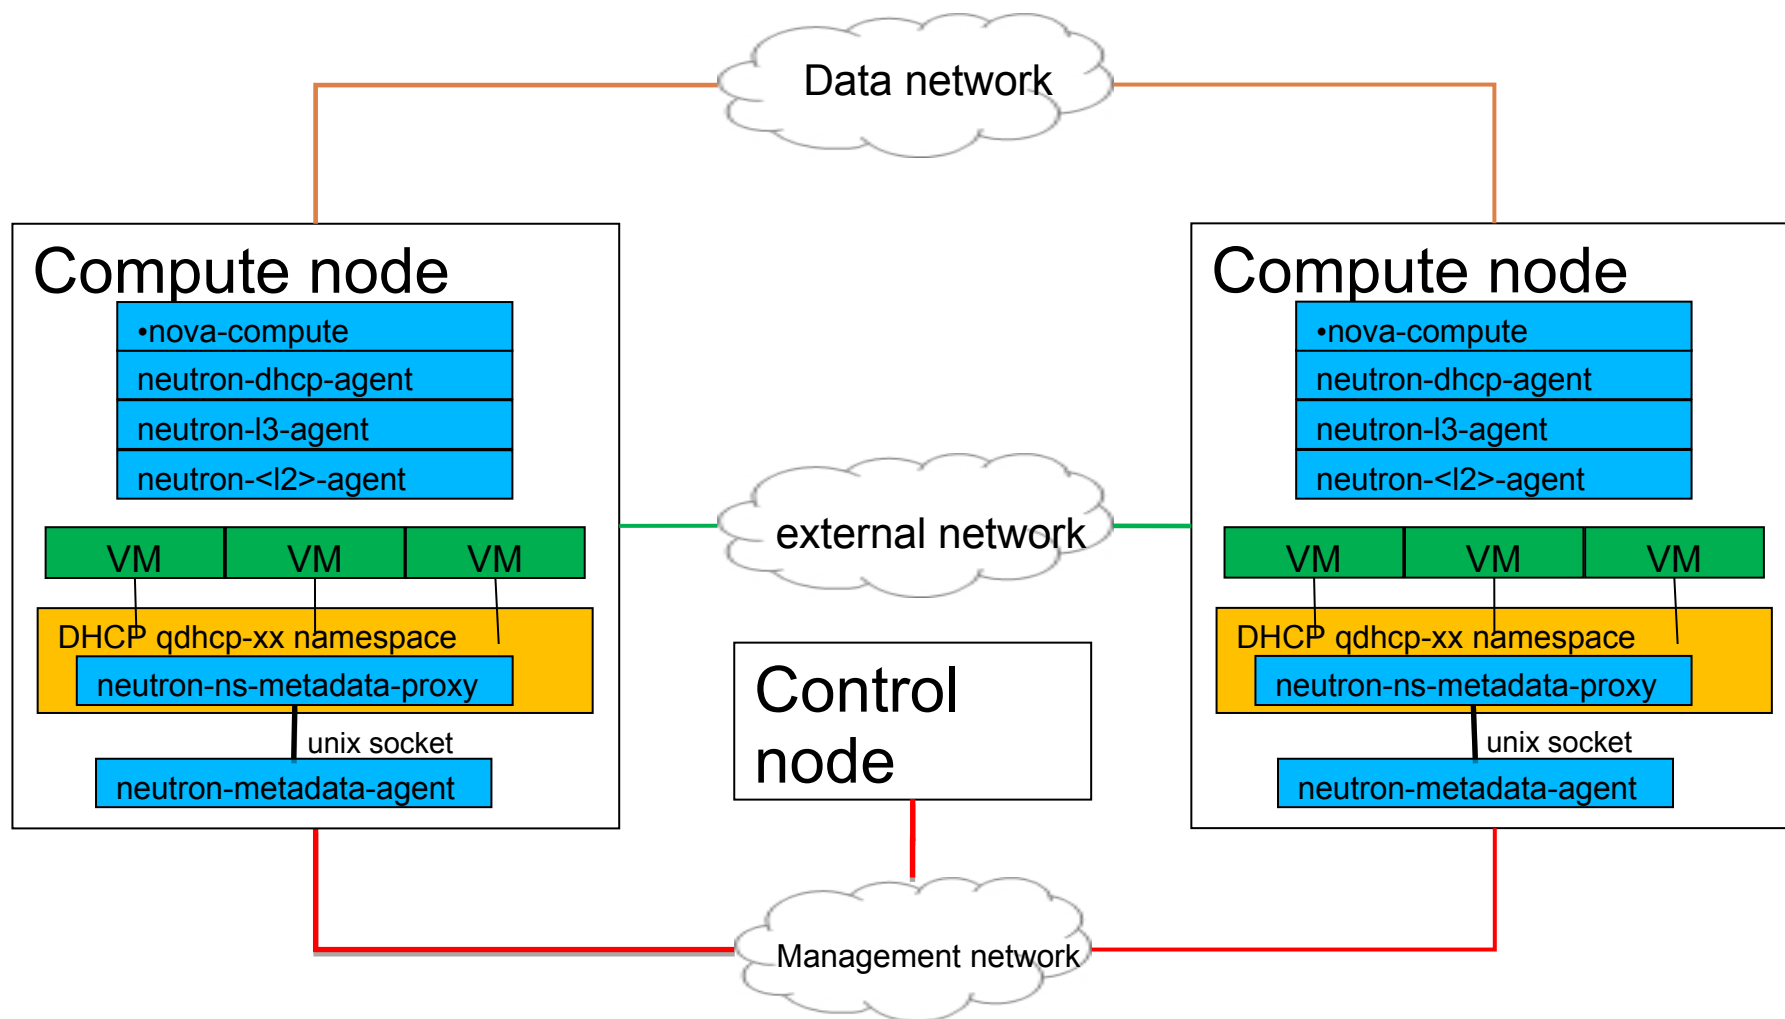
But I have to
duplicate a lot of
DB, segmentation,
etc. work.

What a pain. :(

ML2



Multihost



Thank You



Backup

- 深入理解OpenStack中的网络实现
- <https://github.com/yeasy/easyOVS>